

# Running GGUF Models Offline on Device in Flutter (Tier-1 & Tier-2)


## Overview

This guide explains how to run your fine-tuned Qwen-2.5 GGUF models **locally and offline** inside a Flutter app. The architecture uses two separate models:


- **Tier-1 (Router):** A small, always-on model (Qwen-2.5-1.5B) that classifies intent, routes requests, and outputs strict JSON. It never generates long text.
- **Tier-2 (Reasoner):** A larger, on-demand model (Qwen-2.5-3B or 7B) that handles deep reasoning, explanations, and planning. It outputs natural language.

Both models are exported as .gguf files after fine-tuning and run on-device using a C++ inference engine (llama.cpp) via Flutter's FFI bridge.

## Architecture Overview



## Model Pipeline: Fine-Tune to Deployment



# Understanding GGUF File Formats: F16 vs Q4\_K\_M

---

After fine-tuning, you export your model to GGUF. The two most common formats are **F16** (full precision) and **Q4\_K\_M** (quantized). Understanding the difference is critical for choosing the right file for each tier.

## F16 (Float16 / Half Precision)

- **What it is:** The model weights are stored in 16-bit floating point. This is essentially the full-quality model with no compression.
- **File size:** Large. A 1.5B parameter model is ~3 GB; a 3B model is ~6 GB; a 7B model is ~14 GB.
- **Quality:** Maximum accuracy. No quality loss from the trained weights.
- **Speed:** Slower inference on mobile because more data must be read from memory per token.
- **RAM usage:** High. The entire model must fit in RAM at full precision.
- **When to use F16:**
  - During development and testing on a desktop/emulator to validate model correctness.
  - As the source file for producing quantized versions (you quantize FROM F16).
  - On high-end devices with abundant RAM where maximum quality and speed is not a concern.
  - **Never ship F16 to production mobile apps** unless targeting tablets or very high-end devices with 12+ GB RAM.

## Q4\_K\_M (4-bit Quantization, K-Quant Medium)

- **What it is:** The model weights are compressed to ~4.5 bits per weight using the K-quant method with medium quality settings. This is a lossy compression, but the quality loss is minimal for well-trained models.
- **File size:** Much smaller. A 1.5B parameter model is ~900 MB; a 3B model is ~1.8 GB; a 7B model is ~4 GB.
- **Quality:** Very close to F16 for most tasks. Routing/classification (Tier-1) and natural language generation (Tier-2) both survive Q4\_K\_M quantization well. Minor degradation is possible on edge-case reasoning tasks.
- **Speed:** Faster inference on mobile because less memory bandwidth is needed per token.
- **RAM usage:** Significantly lower. This is what makes on-device inference practical.
- **When to use Q4\_K\_M:**
  - **Production mobile apps** -- this is the standard format for shipping.
  - Both Tier-1 and Tier-2 models in your Flutter app.
  - Any device with limited RAM (which is every phone).

## Quick Comparison Table

| Aspect          | F16 (Half Precision)                 | Q4_K_M (4-bit Quantized)     |
|-----------------|--------------------------------------|------------------------------|
| Bits per weight | 16                                   | ~4.5                         |
| 1.5B model size | ~3 GB                                | ~900 MB                      |
| 3B model size   | ~6 GB                                | ~1.8 GB                      |
| 7B model size   | ~14 GB                               | ~4 GB                        |
| Quality loss    | None                                 | Minimal                      |
| Inference speed | Slower                               | Faster                       |
| RAM required    | High                                 | Low                          |
| Use case        | Dev/testing, source for quantization | Production mobile deployment |

## How to Convert F16 to Q4\_K\_M

If you have an F16 GGUF and need to quantize it:

```
# Using Llama.cpp's quantize tool
./quantize model_f16.gguf model_q4_k_m.gguf Q4_K_M
```

You always fine-tune first, export to F16 GGUF, then quantize to Q4\_K\_M for deployment.

## Which GGUF Files You Need

For the Tier-1 / Tier-2 architecture, you will have **two separate GGUF files**:

| File                     | Model                      | Role                        | Recommended Format | Approximate Size |
|--------------------------|----------------------------|-----------------------------|--------------------|------------------|
| tier1_router_q4_k_m.gguf | Qwen-2.5-1.5B (fine-tuned) | Intent routing, JSON output | Q4_K_M             | ~900 MB          |

---

|                            |                                     |                                   |        |         |
|----------------------------|-------------------------------------|-----------------------------------|--------|---------|
| tier2_reasoner_q4_k_m.gguf | Qwen-<br>2.5-3B<br>(fine-<br>tuned) | Reasoning,<br>natural<br>language | Q4_K_M | ~1.8 GB |
|----------------------------|-------------------------------------|-----------------------------------|--------|---------|

**Total on-device storage:** ~2.7 GB (both models combined with Q4\_K\_M).

**Note:** You do NOT need to load both models into RAM simultaneously. Tier-1 stays loaded. Tier-2 is loaded on-demand and unloaded when not needed.

---

## Step 1: Add Dependencies

---

Update your pubspec.yaml:

```
dependencies:  
  flutter:  
    sdk: flutter  
  llama_cpp_dart: ^0.1.3      # C++ inference engine FFI binding:  
  path_provider: ^2.1.2        # Access to device file system path  
  
flutter:  
  assets:  
    - assets/models/tier1_router_q4_k_m.gguf  
    # Tier-2 is NOT bundled as an asset (too large).  
    # It is downloaded separately after install. See Step 4.
```

### Why Tier-1 is bundled but Tier-2 is not

- **Tier-1 (~900 MB):** Small enough to bundle inside the APK/IPA. The app works immediately after install because the router is always available.
- **Tier-2 (~1.8 GB+):** Too large to bundle. It would bloat the app download beyond acceptable limits. Instead, download it on first launch (or on-demand) to the device's local storage.

If your Tier-1 model is also too large to bundle (e.g., you used a bigger base model), you can download both models post-install and skip the assets declaration entirely.

---

## Step 2: Configure Native Platforms

---

### Android

android/app/src/main/AndroidManifest.xml:

```
<application
    android:largeHeap="true"
    android:requestLegacyExternalStorage="true">
    <!-- android:LargeHeap gives the app access to more RAM -->
    <!-- Required for Loading GGUF models into memory -->
</application>
```

#### android/app/build.gradle:

```
android {
    defaultConfig {
        // Set minimum SDK to 24 (Android 7.0) for 64-bit support
        minSdkVersion 24
    }

    // Prevent compression of GGUF files in the APK
    aaptOptions {
        noCompress 'gguf'
    }
}
```

**Important:** Older Android devices (pre-Android 10) have a **4 GB file size limit** for individual assets. Q4\_K\_M for 1.5B and 3B models are both under this limit.

## iOS

#### ios/Runner/Info.plist:

- Set deployment target to **iOS 14.0+**
- No additional configuration needed for file access since `getApplicationDocumentsDirectory()` is within the app sandbox.

#### ios/Podfile:

```
platform :ios, '14.0'
```

---

## Step 3: Implement Model Services

---

You need two service classes -- one for each tier. Both use the same llama.cpp engine but with different models and different usage patterns.

### Tier-1 Model Service (Always-On Router)

Create lib/services/tier1\_service.dart:



```

import 'dart:io';
import 'dart:convert';
import 'package:flutter/services.dart';
import 'package:path_provider/path_provider.dart';
import 'package:llama_cpp_dart/llama_cpp_dart.dart';

class Tier1Service {
  LlamaProcessor? _processor;
  bool _isLoaded = false;

  /// Copies the Tier-1 model from app assets to local storage.
  /// The C++ engine needs a real file path -- it cannot read from
  /// Flutter's asset bundle directly.
  Future<String> _copyAssetToLocal() async {
    final dir = await getApplicationDocumentsDirectory();
    final filePath = '${dir.path}/tier1_router_q4_k_m.gguf';
    final file = File(filePath);

    if (!await file.exists()) {
      final byteData = await rootBundle.load(
        'assets/models/tier1_router_q4_k_m.gguf',
      );
      await file.writeAsBytes(
        byteData.buffer.asUint8List(
          byteData.offsetInBytes,
          byteData.lengthInBytes,
        ),
      );
    }
    return filePath;
  }

  /// Load Tier-1 into memory. Call this once at app startup.
  /// Tier-1 stays loaded for the entire app session.
  Future<void> load() async {
    if (_isLoaded) return;

    final modelPath = await _copyAssetToLocal();
    final params = ModelParams()
      ..nCtx = 512           // Small context window (routing needs to be fast)
      ..nThreads = 4;       // Adjust based on device CPU cores

    _processor = LlamaProcessor(path: modelPath, modelParams: params);
    _processor?.load();
    _isLoaded = true;
  }

  /// Run a user message through Tier-1 and parse the JSON response.
  /// Returns a structured routing decision.
  Future<Map<String, dynamic>> route(String userMessage) async {
    if (!_isLoaded) await load();
  }
}

```

```

    final systemPrompt = 'You are a Tier-1 router. Output JSON <|im_start|>system\n$systemPrompt<|im_end|>';
    final prompt = '<|im_start|>system\n$systemPrompt<|im_end|>\n' +
        '<|im_start|>user\n$userMessage<|im_end|>\n' +
        '<|im_start|>assistant\n';

    final buffer = StringBuffer();
    await for (final token in _processor!.stream(prompt)) {
        buffer.write(token);
    }

    return jsonDecode(buffer.toString()) as Map<String, dynamic>;
}

void dispose() {
    _processor?.unload();
    _isLoading = false;
}
}

```

### Key points about Tier-1:

- Loaded once at app startup, stays in memory the entire session.
  - Uses a small context window (`nCtx = 512`) because routing decisions are short.
  - Output is always strict JSON matching the Tier-1 schema (`intent`, `journey`, `tool`, `complexity_score`, etc.).
  - Uses **ChatML format** (`<|im_start|>` / `<|im_end|>` tokens) because the Qwen-2.5 models are pretrained with this format.
- 

## Tier-2 Model Service (On-Demand Reasoner)

Create `lib/services/tier2_service.dart`:



```

import 'dart:io';
import 'package:path_provider/path_provider.dart';
import 'package:llama_cpp_dart/llama_cpp_dart.dart';

class Tier2Service {
    LlamaProcessor? _processor;
    bool _isLoaded = false;
    bool _isDownloaded = false;

    /// Check if the Tier-2 model has been downloaded to Local storage.
    Future<bool> isAvailable() async {
        final dir = await getApplicationDocumentsDirectory();
        final file = File('${dir.path}/tier2_reasoner_q4_k_m.gguf');
        _isDownloaded = await file.exists();
        return _isDownloaded;
    }

    /// Download the Tier-2 model from your server.
    /// Call this on first launch or when the user opts in.
    /// Returns a stream of download progress (0.0 to 1.0).
    Stream<double> download(String downloadUrl) async* {
        final dir = await getApplicationDocumentsDirectory();
        final filePath = '${dir.path}/tier2_reasoner_q4_k_m.gguf';
        final file = File(filePath);

        // Use your preferred HTTP client (dio, http, etc.)
        // This is pseudocode -- replace with your actual download logic.
        final httpClient = HttpClient();
        final request = await httpClient.getUrl(Uri.parse(downloadUrl));
        final response = await request.close();
        final totalBytes = response.contentLength;
        var receivedBytes = 0;

        final sink = file.openWrite();
        await for (final chunk in response) {
            sink.add(chunk);
            receivedBytes += chunk.length;
            if (totalBytes > 0) {
                yield receivedBytes / totalBytes;
            }
        }
        await sink.close();
        httpClient.close();
        _isDownloaded = true;
    }

    /// Load Tier-2 into memory. Only call when needed.
    Future<void> ensureLoaded() async {
        if (_isLoaded) return;
        if (!_isDownloaded) {
            final available = await isAvailable();
            if (!available) {

```

```

        throw StateError('Tier-2 model not downloaded yet');
    }
}

final dir = await getApplicationDocumentsDirectory();
final modelPath = '${dir.path}/tier2_reasoner_q4_k_m.gguf';

final params = ModelParams()
..nCtx = 2048           // Larger context for reasoning tasks
..nThreads = 4;

_processor = LlamaProcessor(path: modelPath, modelParams: params);
_processor?.load();
_isLoaded = true;
}

/// Generate a reasoning response. Returns a stream of tokens
/// so you can display the response as it generates (streaming)
Stream<String> generate(String userMessage) {
if (!_isLoaded) {
return Stream.value('Tier-2 model is not loaded.');
}

final systemPrompt = 'You are a helpful reasoning assistant';
final prompt = '<|im_start|>system\n$systemPrompt<|im_end|>\n'
'<|im_start|>user\n$userMessage<|im_end|>\n'
'<|im_start|>assistant\n';

return _processor!.stream(prompt);
}

/// Unload Tier-2 from memory to free RAM.
/// Call this when reasoning is complete or on memory pressure.
void unload() {
_processor?.unload();
_processor = null;
_isLoaded = false;
}

void dispose() => unload();
}

```

### Key points about Tier-2:

- **Not bundled** with the app. Downloaded separately to device storage.
- **Loaded on-demand** only when Tier-1's complexity\_score exceeds the threshold (e.g., > 60).
- **Unloaded after use** to free RAM for the rest of the app.
- Uses a larger context window (nCtx = 2048) because reasoning tasks produce longer outputs.

- Output is natural language, not JSON.
- 

## Step 4: Implement the Orchestrator

---

The orchestrator ties Tier-1 and Tier-2 together. This is the core control logic of your app.

Create `lib/services/orchestrator.dart`:



```

import 'tier1_service.dart';
import 'tier2_service.dart';

class Orchestrator {
    final Tier1Service _tier1 = Tier1Service();
    final Tier2Service _tier2 = Tier2Service();

    static const int complexityThreshold = 60;

    /// Initialize -- Load Tier-1 at startup. Tier-2 stays dormant.
    Future<void> init() async {
        await _tier1.load();
    }

    /// Main entry point: process any user message.
    Future<OrchestratorResult> handle(String userMessage) async {
        // Step 1: Always run Tier-1 first
        final routing = await _tier1.route(userMessage);

        // Step 2: If Tier-1 needs clarification, ask the user
        if (routing['needs_clarification'] == true) {
            return OrchestratorResult(
                type: ResultType.clarification,
                text: routing['clarification'] as String,
            );
        }

        // Step 3: If complexity is high, escalate to Tier-2
        final complexity = routing['complexity_score'] as int? ?? 0;
        if (complexity > complexityThreshold) {
            await _tier2.ensureLoaded();

            final buffer = StringBuffer();
            await for (final token in _tier2.generate(userMessage))
                buffer.write(token);
        }

        // Optionally unload Tier-2 to free memory
        _tier2.unload();

        return OrchestratorResult(
            type: ResultType.reasoning,
            text: buffer.toString(),
            routing: routing,
        );
    }

    // Step 4: If a cloud tool is needed, call it
    if (routing['needs_tools'] == true) {
        return OrchestratorResult(
            type: ResultType.toolCall,
            text: null,
        );
    }
}

```

```

        routing: routing,
        // Flutter will now call the cloud API using
        // routing['tool'] and routing['arguments']
    );
}

// Step 5: Simple response -- Tier-1 handled it entirely
return OrchestratorResult(
    type: ResultType.local,
    routing: routing,
);
}

void dispose() {
    _tier1.dispose();
    _tier2.dispose();
}
}

enum ResultType { clarification, reasoning, toolCall, local }

class OrchestratorResult {
    final ResultType type;
    final String? text;
    final Map<String, dynamic>? routing;

    OrchestratorResult({required this.type, this.text, this.routing});
}

```

---


## Step 5: Memory Management Strategy

Running models on-device requires careful memory management. Here is how to handle both tiers without crashing the app.

### RAM Budget by Device Tier

| Device Class            | Available RAM   | Strategy   |
|-------------------------|-----------------|--|
| Low-end (3-4 GB RAM)    | ~1.5 GB for app | Tier-1 only (Q4_K_M 1.5B). Skip Tier-2.  |
| Mid-range (6-8 GB RAM)  | ~3 GB for app   | Tier-1 always loaded + Tier-2 (3B Q4_K_M) on-demand. Unload Tier-2 after each use. |
| High-end (8-12+ GB RAM) | ~5 GB for app   | Tier-1 always loaded + Tier-2 (3B or 7B Q4_K_M) can stay resident longer.          |

# Memory Lifecycle Diagram



## Memory Rules

1. **Tier-1 stays loaded** for the entire session. At ~900 MB (Q4\_K\_M 1.5B), this is manageable on all modern phones.
2. **Tier-2 loads only when `complexity_score > threshold`** and unloads immediately after the response is complete.
3. **Monitor memory pressure.** On Android, listen to `ComponentCallbacks2.onTrimMemory()`. On iOS, respond to `UIApplication.didReceiveMemoryWarningNotification`. If triggered, unload Tier-2 immediately.
4. **Never load both models at full context simultaneously** on low-end devices. If Tier-2 needs to load, ensure Tier-1's context is minimal (it should be, since routing is already done).

## Step 6: File Structure

---

Here is the recommended project structure for the model-related files:

```
lib/
  services/
    tier1_service.dart      # Tier-1 router model service
    tier2_service.dart      # Tier-2 reasoner model service
    orchestrator.dart       # Coordinates Tier-1, Tier-2, and cloud
  assets/
    models/
      tier1_router_q4_k_m.gguf # Bundled with the app (if small enough)
```

The Tier-2 model file (`tier2_reasoner_q4_k_m.gguf`) lives in the device's application documents directory after download. It is **not** in the `assets/` folder.

---


## Complete Flow: What Happens When a User Sends a Message

---

### Orchestrator Decision Flowchart

This is the core branching logic the orchestrator executes for every incoming user message:






---

## Example 1: High-Complexity Request (Tier-2 Escalation)


**User types:** "Plan a 10-day Europe trip on a budget"



---


## Example 2: Tool Call Request (Cloud API, No Tier-2)

**User types:** "What's the weather in Tokyo?"



### Example 3: Simple Local Response (Tier-1 Only)

**User types:** "Hi, what can you do?"



## Troubleshooting

### Common Issues

| Problem | Cause | Fix |
|---------|-------|-----|
|---------|-------|-----|

|                             |  |   |
|-----------------------------|--|---|
| App crashes on model load   | Not enough RAM                                       | Use Q4_K_M instead of F16. Reduce nCtx. Close background apps during testing.       |
| Model outputs garbage       | Wrong prompt format                                  | Use ChatML format (< im_start > / < im_end > tokens). Qwen-2.5 requires this.       |
| Tier-1 outputs invalid JSON | Poor fine-tuning data or too aggressive quantization | Validate your dataset. Try Q5_K_M if Q4_K_M causes JSON issues.                     |
| Asset copy takes too long   | Large model file                                     | Show a loading screen with progress. The copy only happens once (first launch).     |
| Tier-2 download fails       | Network issues                                       | Implement retry logic and resume support for large file downloads.                  |
| Slow inference              | Too many context tokens or too many threads          | Reduce nCtx. Set nThreads to match physical CPU cores (not logical).                |
| 4 GB asset limit on Android | GGUF file exceeds 4 GB                               | Only affects 7B F16 models. Q4_K_M for 7B is ~4 GB (borderline). Use 3B for safety. |

## Validating Your Model Works

Before integrating into Flutter, test your GGUF file on desktop using the llama.cpp CLI:

```
# Test Tier-1 (should output JSON)
./main -m tier1_router_q4_k_m.gguf \
-p "<|im_start|>system\nYou are a Tier-1 router. Output JSON only"
-n 256

# Test Tier-2 (should output natural language)
./main -m tier2_reasoner_q4_k_m.gguf \
-p "<|im_start|>system\nYou are a helpful reasoning assistant.<|ir"
-n 512
```

---

## Summary: Decision Checklist

---

| Decision                     | Answer                                    |
|------------------------------|---|
| Which format for production? | <b>Q4_K_M</b> for both Tier-1 and Tier-2. |

|                                     |  |
|-------------------------------------|--|
| Which format for development?       | <b>F16</b> to validate quality, then quantize.                   |
| Bundle Tier-1 in the app?           | <b>Yes</b> , if under ~1 GB.                                     |
| Bundle Tier-2 in the app?           | <b>No</b> . Download post-install.                               |
| Load both models at once?           | <b>No</b> . Tier-1 stays loaded. Tier-2 loads/unloads on demand. |
| Which prompt format?                | <b>ChatML</b> (< im_start > / < im_end >).                       |
| Minimum device for Tier-1 only?     | 4 GB RAM (most modern phones).                                   |
| Minimum device for Tier-1 + Tier-2? | 6 GB RAM (mid-range 2022+ phones).                               |