

Running GGUF Models Offline on Device in Flutter (Tier-1 & Tier-2)

Overview

This guide explains how to run your fine-tuned Qwen-2.5 GGUF models **locally and offline** inside a Flutter app. The architecture uses two separate models:

- **Tier-1 (Router):** A small, always-on model (Qwen-2.5-1.5B) that classifies intent, routes requests, and outputs strict JSON. It never generates long text.
- **Tier-2 (Reasoner):** A larger, on-demand model (Qwen-2.5-3B or 7B) that handles deep reasoning, explanations, and planning. It outputs natural language.

Both models are exported as `.gguf` files after fine-tuning and run on-device using a C++ inference engine (`llama.cpp`) via Flutter's FFI bridge.

Understanding GGUF File Formats: F16 vs Q4_K_M

After fine-tuning, you export your model to GGUF. The two most common formats are **F16** (full precision) and **Q4_K_M** (quantized). Understanding the difference is critical for choosing the right file for each tier.

F16 (Float16 / Half Precision)

- **What it is:** The model weights are stored in 16-bit floating point. This is essentially the full-quality model with no compression.
- **File size:** Large. A 1.5B parameter model is ~3 GB; a 3B model is ~6 GB; a 7B model is ~14 GB.
- **Quality:** Maximum accuracy. No quality loss from the trained weights.
- **Speed:** Slower inference on mobile because more data must be read from memory per token.
- **RAM usage:** High. The entire model must fit in RAM at full precision.
- **When to use F16:**
 - During development and testing on a desktop/emulator to validate model correctness.
 - As the source file for producing quantized versions (you quantize FROM F16).
 - On high-end devices with abundant RAM where maximum quality and speed is not a concern.
 - **Never ship F16 to production mobile apps** unless targeting tablets or very high-end devices with 12+ GB RAM.

Q4_K_M (4-bit Quantization, K-Quant Medium)

- **What it is:** The model weights are compressed to ~4.5 bits per weight using the K-quant method with medium quality settings. This is a lossy compression, but the quality loss is minimal for well-trained models.
- **File size:** Much smaller. A 1.5B parameter model is ~900 MB; a 3B model is ~1.8 GB; a 7B model is ~4 GB.
- **Quality:** Very close to F16 for most tasks. Routing/classification (Tier-1) and natural language generation (Tier-2) both survive Q4_K_M quantization well. Minor degradation is possible on edge-case reasoning tasks.
- **Speed:** Faster inference on mobile because less memory bandwidth is needed per token.
- **RAM usage:** Significantly lower. This is what makes on-device inference practical.
- **When to use Q4_K_M:**
 - **Production mobile apps** – this is the standard format for shipping.
 - Both Tier-1 and Tier-2 models in your Flutter app.
 - Any device with limited RAM (which is every phone).

Quick Comparison Table

Aspect	F16 (Half Precision)	Q4_K_M (4-bit Quantized)
Bits per weight	16	~4.5
1.5B model size	~3 GB	~900 MB
3B model size	~6 GB	~1.8 GB
7B model size	~14 GB	~4 GB
Quality loss	None	Minimal
Inference speed	Slower	Faster
RAM required	High	Low
Use case	Dev/testing, source for quantization	Production mobile deployment

How to Convert F16 to Q4_K_M

If you have an F16 GGUF and need to quantize it:

```
# Using llama.cpp's quantize tool
./quantize model_f16.gguf model_q4_k_m.gguf Q4_K_M
```

You always fine-tune first, export to F16 GGUF, then quantize to Q4_K_M for deployment.

Which GGUF Files You Need

For the Tier-1 / Tier-2 architecture, you will have **two separate GGUF files**:

File	Model	Role	Recommended Format	Approximate Size
tier1_router_q4_k_m.gguf	Qwen-2.5-1.5B (fine-tuned)	Intent routing, JSON output	Q4_K_M	~900 MB
tier2_reasoner_q4_k_m.gguf	Qwen-2.5-3B (fine-tuned)	Reasoning, natural language	Q4_K_M	~1.8 GB

Total on-device storage: ~2.7 GB (both models combined with Q4_K_M).

Note: You do NOT need to load both models into RAM simultaneously. Tier-1 stays loaded. Tier-2 is loaded on-demand and unloaded when not needed.

Step 1: Add Dependencies

Update your `pubspec.yaml`:

```
dependencies:
  flutter:
    sdk: flutter
  llama_cpp_dart: ^0.1.3      # C++ inference engine FFI bindings
  path_provider: ^2.1.2        # Access to device file system paths

flutter:
  assets:
    - assets/models/tier1_router_q4_k_m.gguf
    # Tier-2 is NOT bundled as an asset (too large).
    # It is downloaded separately after install. See Step 4.
```

Why Tier-1 is bundled but Tier-2 is not

- **Tier-1 (~900 MB):** Small enough to bundle inside the APK/IPA. The app works immediately after install because the router is always available.
- **Tier-2 (~1.8 GB+):** Too large to bundle. It would bloat the app download beyond acceptable limits. Instead, download it on first launch (or on-demand) to the device's local storage.

If your Tier-1 model is also too large to bundle (e.g., you used a bigger base model), you can download both models post-install and skip the `assets` declaration entirely.

Step 2: Configure Native Platforms

Android

`android/app/src/main/AndroidManifest.xml`:

```
<application
  android:largeHeap="true"
  android:requestLegacyExternalStorage="true">
  <!-- android:largeHeap gives the app access to more RAM -->
  <!-- Required for loading GGUF models into memory -->
</application>
```

`android/app/build.gradle`:

```

    android {
        defaultConfig {
            // Set minimum SDK to 24 (Android 7.0) for 64-bit support
            minSdkVersion 24
        }

        // Prevent compression of GGUF files in the APK
        aaptOptions {
            noCompress 'gguf'
        }
    }
}

```

Important: Older Android devices (pre-Android 10) have a **4 GB file size limit** for individual assets. Q4_K_M for 1.5B and 3B models are both under this limit.

iOS

ios/Runner/Info.plist:

- Set deployment target to **iOS 14.0+**
- No additional configuration needed for file access since `getApplicationDocumentsDirectory()` is within the app sandbox.

ios/Podfile:

```
platform :ios, '14.0'
```

Step 3: Implement Model Services

You need two service classes – one for each tier. Both use the same `llama.cpp` engine but with different models and different usage patterns.

Tier-1 Model Service (Always-On Router)

Create `lib/services/tier1_service.dart`:

```

import 'dart:io';
import 'dart:convert';
import 'package:flutter/services.dart';
import 'package:path_provider/path_provider.dart';
import 'package:llama_cpp_dart/llama_cpp_dart.dart';

class Tier1Service {
    LlamaProcessor? _processor;
    bool _isLoaded = false;

    /// Copies the Tier-1 model from app assets to local storage.
    /// The C++ engine needs a real file path -- it cannot read from
    /// Flutter's asset bundle directly.
    Future<String> _copyAssetToLocal() async {
        final dir = await getApplicationDocumentsDirectory();
        final filePath = '${dir.path}/tier1_router_q4_k_m.gguf';
        final file = File(filePath);

        if (!await file.exists()) {
            final byteData = await rootBundle.load(
                'assets/models/tier1_router_q4_k_m.gguf',
            );
            await file.writeAsBytes(
                byteData.buffer.asUint8List(
                    byteData.offsetInBytes,
                    byteData.lengthInBytes,
                ),
            );
        }
        return filePath;
    }
}

```

```

/// Load Tier-1 into memory. Call this once at app startup.
/// Tier-1 stays loaded for the entire app session.
Future<void> load() async {
    if (_isLoaded) return;

    final modelPath = await _copyAssetToLocal();
    final params = ModelParams()
        ..nCtx = 512          // Small context window (routing needs little)
        ..nThreads = 4;      // Adjust based on device CPU cores

    _processor = LlamaProcessor(path: modelPath, modelParams: params);
    _processor?.load();
    _isLoaded = true;
}

/// Run a user message through Tier-1 and parse the JSON response.
/// Returns a structured routing decision.
Future<Map<String, dynamic>> route(String userMessage) async {
    if (!_isLoaded) await load();

    final systemPrompt = 'You are a Tier-1 router. Output JSON only.';
    final prompt = '<|im_start|>system\n$systemPrompt<|im_end|>\n'
        '<|im_start|>user\n$userMessage<|im_end|>\n'
        '<|im_start|>assistant\n';

    final buffer = StringBuffer();
    await for (final token in _processor!.stream(prompt)) {
        buffer.write(token);
    }

    return jsonDecode(buffer.toString()) as Map<String, dynamic>;
}

void dispose() {
    _processor?.unload();
    _isLoaded = false;
}
}

```

Key points about Tier-1:

- Loaded once at app startup, stays in memory the entire session.
- Uses a small context window (nCtx = 512) because routing decisions are short.
- Output is always strict JSON matching the Tier-1 schema (intent, journey, tool, complexity_score, etc.).
- Uses ChatML format (<|im_start|> / <|im_end|> tokens) because the Qwen-2.5 models are pretrained with this format.

Tier-2 Model Service (On-Demand Reasoner)

Create lib/services/tier2_service.dart:

```

import 'dart:io';
import 'package:path_provider/path_provider.dart';
import 'package:llama_cpp_dart/llama_cpp_dart.dart';

class Tier2Service {
    LlamaProcessor? _processor;
    bool _isLoaded = false;
    bool _isDownloaded = false;

    /// Check if the Tier-2 model has been downloaded to local storage.
    Future<bool> isAvailable() async {
        final dir = await getApplicationDocumentsDirectory();
        final file = File('${dir.path}/tier2_reasoner_q4_k_m.gguf');
        _isDownloaded = await file.exists();
    }
}

```

```

        return _isDownloaded;
    }

/// Download the Tier-2 model from your server.
/// Call this on first launch or when the user opts in.
/// Returns a stream of download progress (0.0 to 1.0).
Stream<double> download(String downloadUrl) async* {
    final dir = await getApplicationDocumentsDirectory();
    final filePath = '${dir.path}/tier2_reasoner_q4_k_m.gguf';
    final file = File(filePath);

    // Use your preferred HTTP client (dio, http, etc.)
    // This is pseudocode -- replace with your actual download logic.
    final httpClient = HttpClient();
    final request = await httpClient.getUrl(Uri.parse(downloadUrl));
    final response = await request.close();
    final totalBytes = response.contentLength;
    var receivedBytes = 0;

    final sink = file.openWrite();
    await for (final chunk in response) {
        sink.add(chunk);
        receivedBytes += chunk.length;
        if (totalBytes > 0) {
            yield receivedBytes / totalBytes;
        }
    }
    await sink.close();
    httpClient.close();
    _isDownloaded = true;
}

/// Load Tier-2 into memory. Only call when needed.
Future<void> ensureLoaded() async {
    if (_isLoaded) return;
    if (!_isDownloaded) {
        final available = await isAvailable();
        if (!available) {
            throw StateError('Tier-2 model not downloaded yet');
        }
    }

    final dir = await getApplicationDocumentsDirectory();
    final modelPath = '${dir.path}/tier2_reasoner_q4_k_m.gguf';

    final params = ModelParams()
        ..nCtx = 2048      // Larger context for reasoning tasks
        ..nThreads = 4;

    _processor = LlamaProcessor(path: modelPath, modelParams: params);
    _processor?.load();
    _isLoaded = true;
}

/// Generate a reasoning response. Returns a stream of tokens
/// so you can display the response as it generates (streaming UX).
Stream<String> generate(String userMessage) {
    if (!_isLoaded) {
        return Stream.value('Tier-2 model is not loaded.');
    }

    final systemPrompt = 'You are a helpful reasoning assistant.';
    final prompt = '<|im_start|>system\n$systemPrompt<|im_end|>\n'
        '<|im_start|>user\n$userMessage<|im_end|>\n'
        '<|im_start|>assistant\n';
}

```

```

        return _processor!.stream(prompt);
    }

    /// Unload Tier-2 from memory to free RAM.
    /// Call this when reasoning is complete or on memory pressure.
    void unload() {
        _processor?.unload();
        _processor = null;
        _isLoaded = false;
    }

    void dispose() => unload();
}

```

Key points about Tier-2:

- **Not bundled** with the app. Downloaded separately to device storage.
- **Loaded on-demand** only when Tier-1's `complexity_score` exceeds the threshold (e.g., > 60).
- **Unloaded after use** to free RAM for the rest of the app.
- Uses a larger context window (`nCtx = 2048`) because reasoning tasks produce longer outputs.
- Output is natural language, not JSON.

Step 4: Implement the Orchestrator

The orchestrator ties Tier-1 and Tier-2 together. This is the core control logic of your app.

Create `lib/services/orchestrator.dart`:

```

import 'tier1_service.dart';
import 'tier2_service.dart';

class Orchestrator {
    final Tier1Service _tier1 = Tier1Service();
    final Tier2Service _tier2 = Tier2Service();

    static const int complexityThreshold = 60;

    /// Initialize -- load Tier-1 at startup. Tier-2 stays dormant.
    Future<void> init() async {
        await _tier1.load();
    }

    /// Main entry point: process any user message.
    Future<OrchestratorResult> handle(String userMessage) async {
        // Step 1: Always run Tier-1 first
        final routing = await _tier1.route(userMessage);

        // Step 2: If Tier-1 needs clarification, ask the user
        if (routing['needs_clarification'] == true) {
            return OrchestratorResult(
                type: ResultType.clarification,
                text: routing['clarification'] as String,
            );
        }

        // Step 3: If complexity is high, escalate to Tier-2
        final complexity = routing['complexity_score'] as int? ?? 0;
        if (complexity > complexityThreshold) {
            await _tier2.ensureLoaded();

            final buffer = StringBuffer();
            await for (final token in _tier2.generate(userMessage)) {
                buffer.write(token);
            }
        }
    }
}

```

```

        // Optionally unload Tier-2 to free memory
        _tier2.unload();

        return OrchestratorResult(
            type: ResultType.reasoning,
            text: buffer.toString(),
            routing: routing,
        );
    }

    // Step 4: If a cloud tool is needed, call it
    if (routing['needs_tools'] == true) {
        return OrchestratorResult(
            type: ResultType.toolCall,
            text: null,
            routing: routing,
            // Flutter will now call the cloud API using
            // routing['tool'] and routing['arguments']
        );
    }

    // Step 5: Simple response -- Tier-1 handled it entirely
    return OrchestratorResult(
        type: ResultType.local,
        routing: routing,
    );
}

void dispose() {
    _tier1.dispose();
    _tier2.dispose();
}
}

enum ResultType { clarification, reasoning, toolCall, local }

class OrchestratorResult {
    final ResultType type;
    final String? text;
    final Map<String, dynamic>? routing;

    OrchestratorResult({required this.type, this.text, this.routing});
}

```

Step 5: Memory Management Strategy

Running models on-device requires careful memory management. Here is how to handle both tiers without crashing the app.

RAM Budget by Device Tier

Device Class	Available RAM	Strategy
Low-end (3-4 GB RAM)	~1.5 GB for app	Tier-1 only (Q4_K_M 1.5B). Skip Tier-2.
Mid-range (6-8 GB RAM)	~3 GB for app	Tier-1 always loaded + Tier-2 (3B Q4_K_M) on-demand. Unload Tier-2 after each use.
High-end (8-12+ GB RAM)	~5 GB for app	Tier-1 always loaded + Tier-2 (3B or 7B Q4_K_M) can stay resident longer.

Memory Rules

1. **Tier-1 stays loaded** for the entire session. At ~900 MB (Q4_K_M 1.5B), this is manageable on all modern phones.
2. **Tier-2 loads only when `complexity_score > threshold`** and unloads immediately after the response is complete.
3. **Monitor memory pressure**. On Android, listen to `ComponentCallbacks2.onTrimMemory()`. On iOS, respond to `UIApplication.didReceiveMemoryWarningMemoryWarningNotification`. If triggered, unload Tier-2 immediately.
4. **Never load both models at full context simultaneously** on low-end devices. If Tier-2 needs to load, ensure Tier-1's context is minimal (it should be, since routing is already done).

Step 6: File Structure

Here is the recommended project structure for the model-related files:

```
lib/
  services/
    tier1_service.dart      # Tier-1 router model service
    tier2_service.dart      # Tier-2 reasoner model service
    orchestrator.dart       # Coordinates Tier-1, Tier-2, and cloud tools
  assets/
    models/
      tier1_router_q4_k_m.gguf # Bundled with the app (if small enough)
```

The Tier-2 model file (`tier2_reasoner_q4_k_m.gguf`) lives in the device's application documents directory after download. It is **not** in the `assets/` folder.

Complete Flow: What Happens When a User Sends a Message

```
User types: "Plan a 10-day Europe trip on a budget"
  |
  v
[Tier-1 Router] (always in memory)
Outputs JSON:
{
  "intent": "trip_planning",
  "complexity_score": 85,
  "routing_confidence": 0.92,
  "needs_tools": false,
  "needs_clarification": false
}
  |
  v
Flutter reads complexity_score = 85 (> 60 threshold)
  |
  v
[Load Tier-2 into memory]
  |
  v
[Tier-2 Reasoner] generates detailed trip plan
(natural language, streamed to UI token by token)
  |
  v
[Unload Tier-2 from memory]
  |
  v
Response displayed to user
```

```

User types: "What's the weather in Tokyo?"
|
v
[Tier-1 Router] (always in memory)
Outputs JSON:
{
  "intent": "weather_lookup",
  "complexity_score": 25,
  "needs_tools": true,
  "tool": "weather_api",
  "arguments": { "city": "Tokyo" }
}
|
v
Flutter reads needs_tools = true
Tier-2 NOT loaded (complexity_score = 25)
|
v
[Cloud API call] weather_api({ city: "Tokyo" })
|
v
Format JSON result and display to user

```

Troubleshooting

Common Issues

Problem	Cause	Fix
App crashes on model load	Not enough RAM	Use Q4_K_M instead of F16. Reduce nCtx. Close background apps during testing.
Model outputs garbage	Wrong prompt format	Use ChatML format (< im_start > / < im_end > tokens). Qwen-2.5 requires this.
Tier-1 outputs invalid JSON	Poor fine-tuning data or too aggressive quantization	Validate your dataset. Try Q5_K_M if Q4_K_M causes JSON issues.
Asset copy takes too long	Large model file	Show a loading screen with progress. The copy only happens once (first launch).
Tier-2 download fails	Network issues	Implement retry logic and resume support for large file downloads.
Slow inference	Too many context tokens or too many threads	Reduce nCtx. Set nThreads to match physical CPU cores (not logical).
4 GB asset limit on Android	GGUF file exceeds 4 GB	Only affects 7B F16 models. Q4_K_M for 7B is ~4 GB (borderline). Use 3B for safety.

Validating Your Model Works

Before integrating into Flutter, test your GGUF file on desktop using the llama.cpp CLI:

```

# Test Tier-1 (should output JSON)
./main -m tier1_router_q4_k_m.gguf \
-p "<|im_start|>system\nYou are a Tier-1 router. Output JSON only.<|im_end|>\n<|im_start|>user\nBook a flight to Paris<|im_end|>\n-n 256

# Test Tier-2 (should output natural language)
./main -m tier2_reasoner_q4_k_m.gguf \
-p "<|im_start|>system\nYou are a helpful reasoning assistant.<|im_end|>\n<|im_start|>user\nExplain quantum computing simply<|im_end|>\n-n 512

```

Summary: Decision Checklist

Decision	Answer
Which format for production?	Q4_K_M for both Tier-1 and Tier-2.
Which format for development?	F16 to validate quality, then quantize.
Bundle Tier-1 in the app?	Yes , if under ~1 GB.
Bundle Tier-2 in the app?	No . Download post-install.
Load both models at once?	No . Tier-1 stays loaded. Tier-2 loads/unloads on demand.
Which prompt format?	ChatML (< im_start > / < im_end >).
Minimum device for Tier-1 only?	4 GB RAM (most modern phones).
Minimum device for Tier-1 + Tier-2?	6 GB RAM (mid-range 2022+ phones).